



# Machine Learning I

Data Mining and Data integration in Biomedicine  
Master in Bioinformatics

Janet Piñero  
Medbioinformatics Solutions SL  
2025-2026



## Outline

What is machine learning

Classification

K nearest neighbor

Decision Trees

Hands-on session

References

## What is machine learning



"Field of study that gives computers the ability to learn without being explicitly programmed"  
Arthur Samuel (1959)

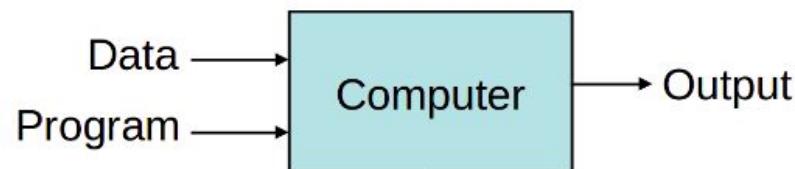
a subset of Artificial Intelligence (AI) in the field of computer science that often uses statistical techniques to give computers the ability to "learn" (i.e., progressively improve performance on a specific task) with data, without being explicitly programmed.

A well posed learning problem: "A computer program is said to learn from experience E with respect to some class of tasks T and performance measure P, if its performance at tasks in T, as measured by P, improves with experience E." Tom Michel (1999)

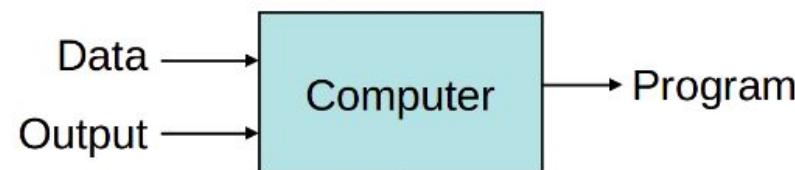
# What is machine learning



## Traditional Programming



## Machine Learning

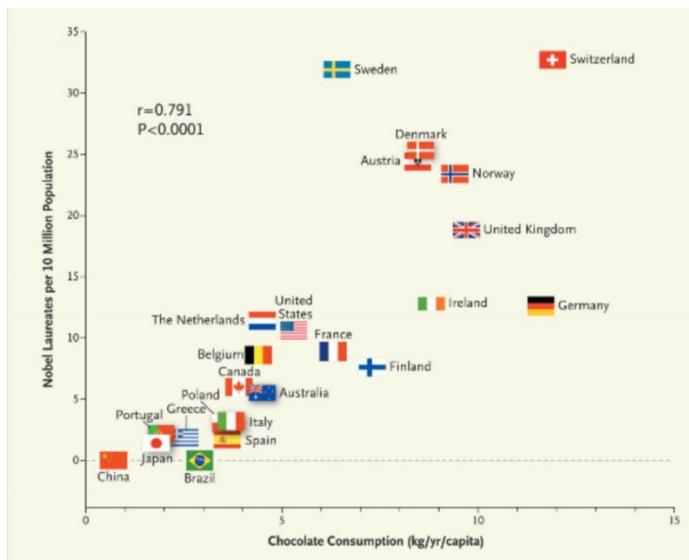


## Components of a predictor



question -> input data -> features -> algorithm -> parameters -> evaluation

# Unrelated data is the most common mistake

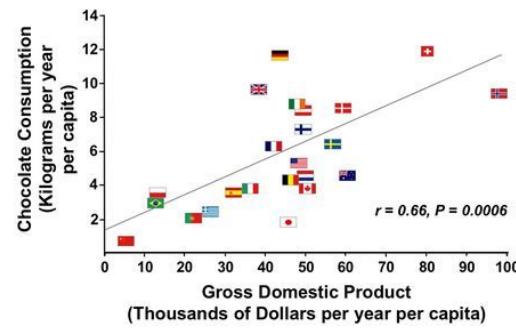
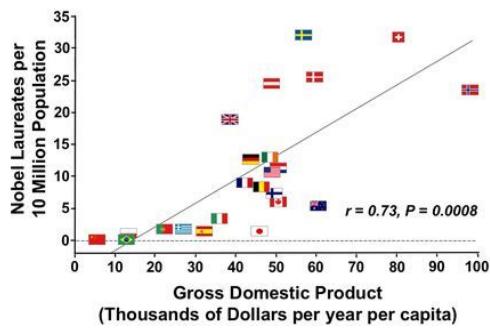
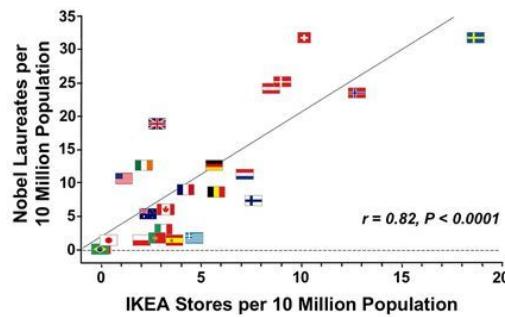


- There was a close, significant linear correlation ( $r=0.791$ ,  $P<0.0001$ ) between chocolate consumption per capita and the number of Nobel laureates per 10 million persons in a total of 23 countries
- When recalculated with the exclusion of Sweden, the correlation coefficient increased to 0.862.
- Switzerland was the top performer in terms of both the number of Nobel laureates and chocolate consumption.
- The slope of the regression line allows us to estimate that it would take about 0.4 kg of chocolate per capita per year to increase the number of Nobel laureates in a given country by 1.

<http://www.nejm.org/doi/full/10.1056/NEJMOn1211064>

## Unrelated data is the most common mistake

DMI





## Components of a predictor

question -> **input data** -> features -> algorithm -> parameters -> evaluation

Often more data > better models

Data relevant for the question is the most important step!



## Components of a predictor

question -> input data -> **features** -> algorithm -> parameters -> evaluation

### Properties of good features

- Lead to data compression
- Retain relevant information
- Are created based on expert application knowledge

### Common mistakes

- Trying to automate feature selection (**automate with care**)
- Not paying attention to data-specific quirks
- Throwing away information unnecessarily

## Components of a predictor



question -> input data -> features -> **algorithm** -> parameters -> evaluation

TABLE 1  
*Performance of linear discriminant analysis and the best result we found on ten randomly selected data sets*

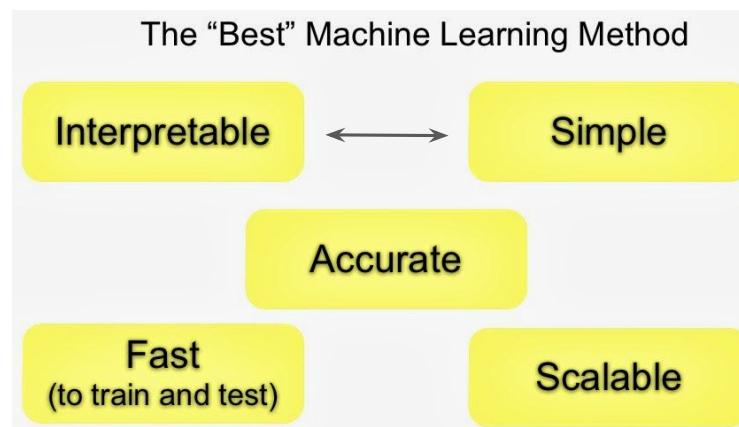
Data set	Best method e.r.	Lindisc e.r.	Default rule	Prop linear
Segmentation	0.0140	0.083	0.760	0.907
Pima	0.1979	0.221	0.350	0.848
House-votes16	0.0270	0.046	0.386	0.948
Vehicle	0.1450	0.216	0.750	0.883
Satimage	0.0850	0.160	0.758	0.889
Heart Cleveland	0.1410	0.141	0.560	1.000
Splice	0.0330	0.057	0.475	0.945
Waveform21	0.0035	0.004	0.667	0.999
Led7	0.2650	0.265	0.900	1.000
Breast Wisconsin	0.0260	0.038	0.345	0.963

e.r. -> error rate  
Default rule -> assigning every point to the majority class

## Components of a predictor



question -> input data -> features -> **algorithm** -> parameters -> evaluation



## **Relative order of importance**



question > data > features > algorithms

## Supervised vs Unsupervised learning



- Supervised Learning:
  - Use a data set X to predict or detect association with a response y
  - Teach the computer how to do something, then let it use its new found knowledge to do it
    - Regression
    - Classification
- Unsupervised Learning:
  - Discover the signal in X; no y is available
  - Let the computer learn how to do something, and use this to determine structure and patterns in data
    - Clustering
    - PCA

# Steps in supervised machine learning



**Pre-processing data:** Data cleaning, normalization and data transformation procedures.

**Training and test data split:** Decide which strategy you want to use for evaluation purposes. You need to use a test set to evaluate your model later on.

**Training the model:** This is where your choice of supervised learning algorithm becomes relevant.

“Training” generally means your data set is used in optimization of the loss function to find parameters for  $f(x)$ .

**Estimating performance of the model:** This is about which metrics to use to evaluate performance and how to calculate those metrics.

**Model tuning and selection:** Trying different parameters and selecting the best model.

# The machine learning process



A major goal of the machine learning process is to find an **algorithm**  $f(X)$  that most accurately predicts future values ( $\hat{Y}$ ) based on a set of features ( $X$ ).

The algorithm not only fits well past data, but more importantly, predicts a future outcome accurately. This is called the **generalizability** of the algorithm.

To assess this, it is crucial how data is divided.

To provide an accurate understanding of the generalizability of the final optimal model, data is split into training and test data sets:

- **Training set:** these data are used to develop feature sets, train algorithms, tune hyperparameters, compare models, and all of the other activities required to choose a final model.
- **Test set:** having chosen a final model, these data are used to estimate an unbiased assessment of the model's performance (*generalization error*).

It is critical that the test set not be used prior to selecting your final model. Assessing results on the test set prior to final model selection biases the model selection process since the testing data will have become part of the model development process.

## Selecting the algorithm

While certain algorithms tend to perform better than others with certain types of data, no single algorithm will always outperform all others on all problems. This concept is called the ***no free lunch theorem***: you don't get something for nothing; you need to put some effort into working out the best algorithm for your particular problem.

Data scientists typically choose a few algorithms they know tend to work well for the type of data and problem they are working on, and see which algorithm generates the best-performing model.

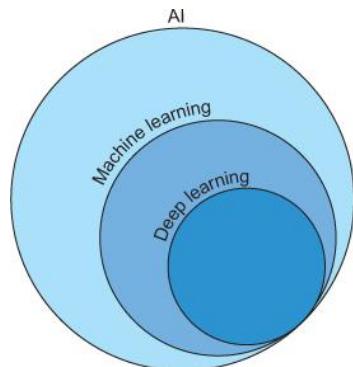
The theorem states that all optimization algorithms perform equally well when their performance is averaged across all possible problems.

## Flavors of Supervised Learning



- Regression: Predict a quantitative response, such as blood pressure, cholesterol level, tumor size
- Classification: Predict a categorical response, such as tumor versus normal tissue, heart disease versus no heart disease, subtype of glioblastoma
- Other Flavors: survival analysis, and more . . .

## Note on deep learning



Very popular in the last 5-10 years:

- models with outstanding performance.
- computational power

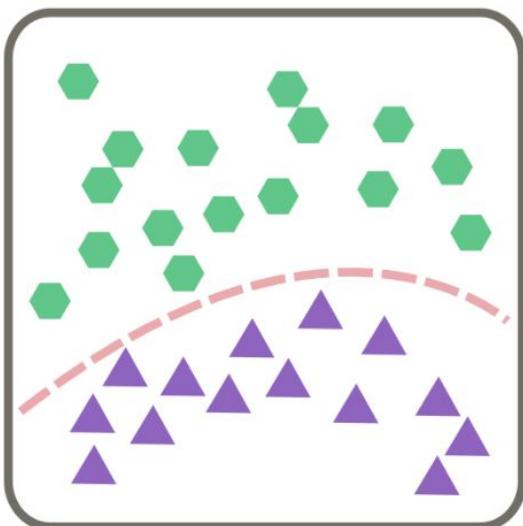
BUT

- Require more data
- Long time (hours or even days!) to train.
- Computationally expensive (a lot of computing power)
- The rules are less interpretable. By their nature, deep learning models favor performance over model interpretability.

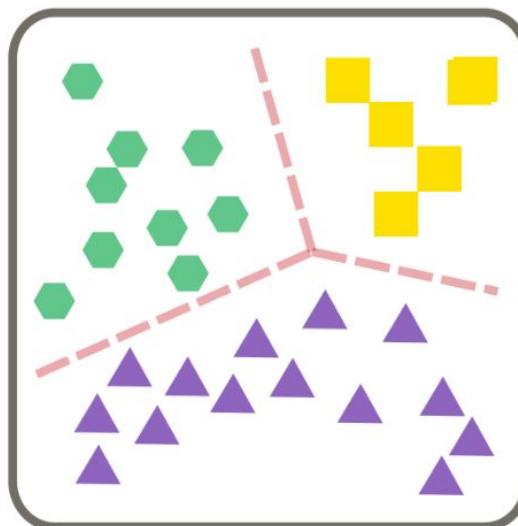
Deep learning algorithms are particularly good at tasks involving complex data, such as image classification, NLP and audio transcription.

# Classification

DMI



Binary classification



Multi-class classification

# Classification vs. Numeric Prediction



## Classification

Predict categorical class labels (discrete or nominal)

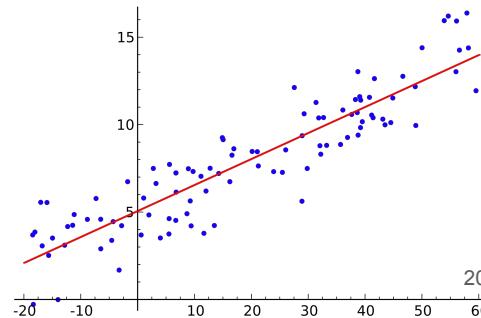
- Construct a model based on the training set and the **class labels** (the values in a classifying attribute) and use it in classifying new data

## Numeric prediction

Model continuous-valued functions (i.e., predict unknown or missing values)

Typical applications of classification

- Medical diagnosis: if a tumor is cancerous or benign
- Disease vs non-disease
- Predicting embryonic cell states



# Classification—Model Construction, Validation and Testing



## Model construction

- Each sample is assumed to belong to a predefined class (shown by the **class label**)
- The set of samples used for model construction is **training set**
- Model: Represented as decision trees, rules, mathematical formulas, or other forms

## Model Validation and Testing:

- **Test:** Estimate accuracy of the model
  - The known label of test sample is compared with the classified result from the model
  - *Accuracy:* % of test set samples that are correctly classified by the model
  - Test set is independent of training set
- **Validation:** If *the test set* is used to select or refine models, it is called **validation** (or development) (**test** set)

**Model Deployment:** If the accuracy is acceptable, use the model to classify new data

## Classification: Definition



- Given a collection of records (training set )
  - Each record contains a set of attributes, one of the attributes is the class.
- Find a model for class attribute as a function of the values of other attributes.
- Goal: previously unseen records should be assigned a class as accurately as possible.
  - A test set is used to determine the accuracy of the model. Usually, the given data set is divided into training and test sets, with training set used to build the model and test set used to validate



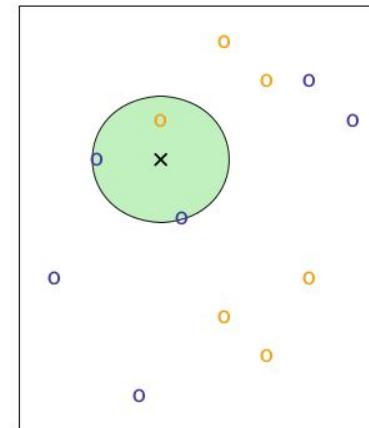
## Classification Techniques

- Logistic Regression
- Decision Tree
- Naïve Bayes
- Rule-based Classification
- K Nearest Neighbor
- Support Vector Machines
- Ensemble methods

## K-Nearest Neighbors

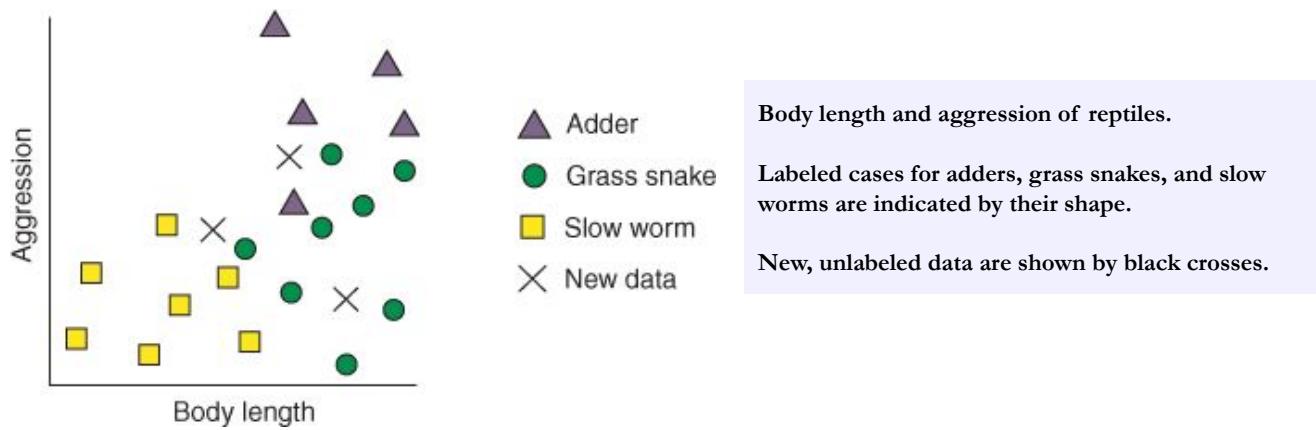


- K-nearest neighbor (KNN) is a very simple algorithm in which each observation is predicted based on its “similarity” to other observations.
- The most common distance measures are the Euclidean and Manhattan distance metrics
- The KNN algorithm identifies  $k$  observations that are “similar” or nearest to the new record being predicted and then uses the average response value (regression) or the most common class (classification) of those  $k$  observations as the predicted output.
- 



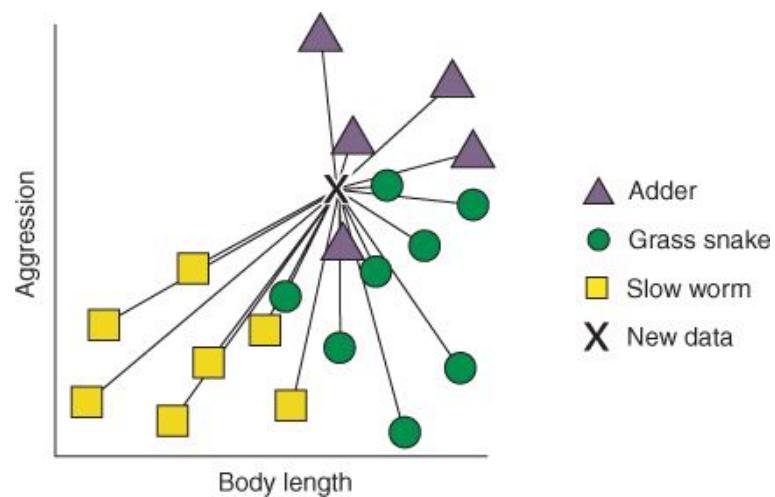
## K-Nearest Neighbors

DMI



## K-Nearest Neighbors

DMI

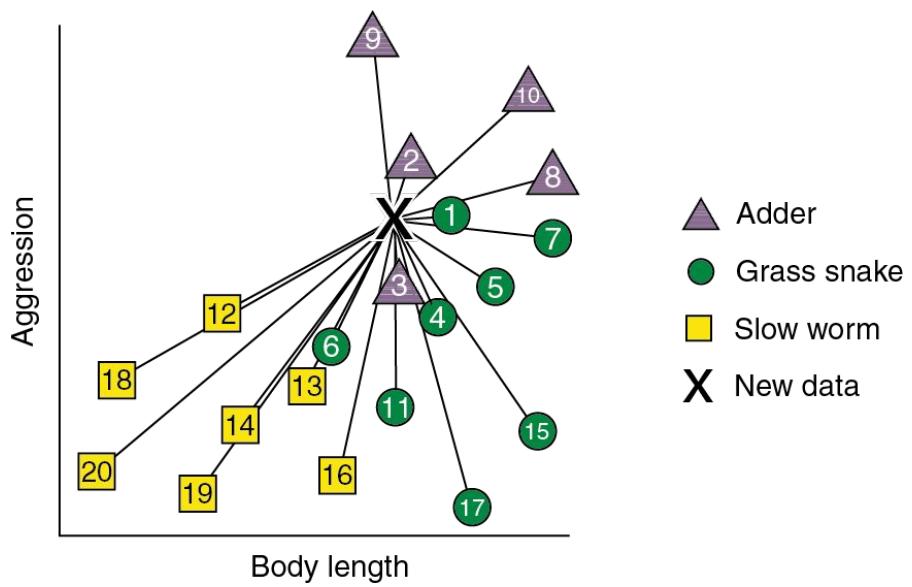


The first step of the kNN algorithm:  
calculating distance.

The lines represent the distance between one  
of the unlabeled cases (the cross) and each of  
the labeled cases.

# K-Nearest Neighbors

DMI

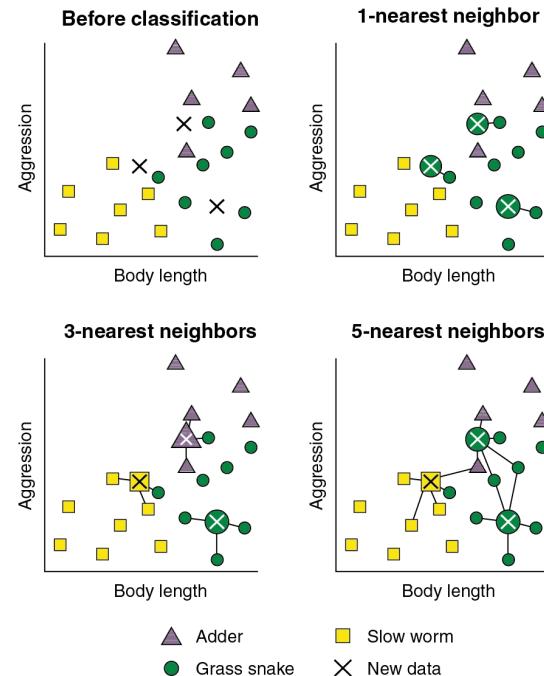


The second step of the kNN algorithm:  
ranking the neighbors.

The lines represent the distance between  
one of the unlabeled cases (the cross)  
and each of the labeled cases.

The numbers represent the ranked  
distance between the unlabeled case (the  
cross) and each labeled case (1 =  
closest).

## K-Nearest Neighbors



The final step of the kNN algorithm:  
identifying the k-nearest neighbors and  
taking the majority vote.

Lines connect the unlabeled data with their  
one, three, and five nearest neighbors.

The majority vote in each scenario is  
indicated by the shape drawn under each  
cross.

used for both classification and  
regression problems

## K-Nearest Neighbors



- ✓ Different features might have different scales. For example, we can have a measure of pain scaling from 1 to 10 or 1 to 100.
- ✓ Some similarity or distances measures assume the same measuring unit in all feature dimensions.
- ✓ This requires that the data may need to be transferred into the same scale.
- ✓ Re-scaling can make each feature contribute to the distance in a relatively equal manner, avoiding potential bias.

## KNNs: Choosing k

The performance of KNNs is very sensitive to the choice of  $k$ :

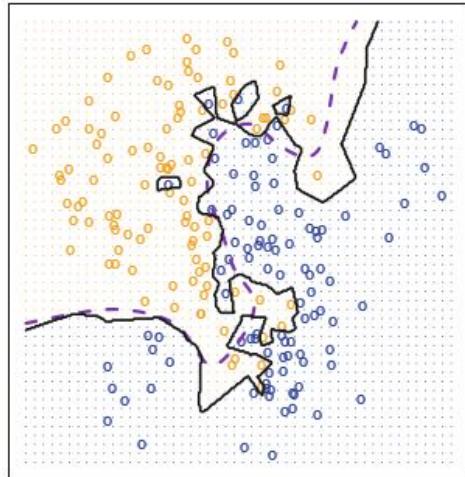
- In general, low values of  $k$  typically overfit and large values often underfit.
- At the extremes, when  $k=1$ , we base our prediction on a single observation that has the closest distance measure
- When  $k=n$ , we are simply using the average (regression) or most common class (classification) across all training samples as our predicted value.

There is no general rule about the best  $k$  as it depends greatly on the nature of the data. For high signal data with very few noisy (irrelevant) features, smaller values of  $k$  tend to work best. As more irrelevant features are involved, larger values of  $k$  are required to smooth out the noise.

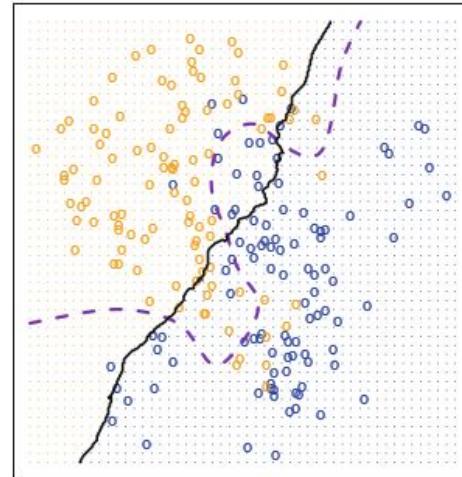
**When using KNN for classification, it is best to assess odd numbers for  $k$  to avoid ties in the event there is equal proportion of response levels**

## KNNs: Choosing k

KNN: K=1



KNN: K=100



A comparison of the KNN decision boundaries (solid black curves) obtained using  $K = 1$  and  $K = 100$

## KNNs



KNNs are a very simplistic, and intuitive, algorithm that can provide average to decent predictive power, especially when the response is dependent on the local structure of the features.

A major drawback of KNNs is their computation time, which increases by  $n \times p$  for each observation.

Although KNNs rarely provide the best predictive performance, they have many benefits, for example, in feature engineering and in data cleaning and preprocessing (missing data imputation)

# Predicting with trees



## Key ideas

- Iteratively split variables into groups
- Evaluate "homogeneity" within each group
- Split again if necessary

## Pros:

- Easy to interpret
- Better performance in nonlinear settings

Tree-based models are a class of nonparametric algorithms that work by partitioning the feature space into a number of smaller (non-overlapping) regions with similar response values using a set of splitting rules.

## Cons:

- Without pruning/cross-validation can lead to overfitting
- Harder to estimate uncertainty
- Results may be variable

## Example: the breast cancer dataset

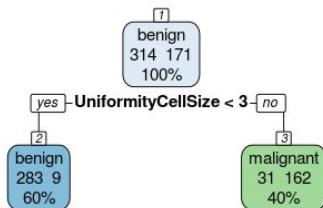


Sample number	Clump Thickness	Uniformity of Cell Size	Uniformity of Cell Shape	Marginal Adhesion	Single Epithelial Cell Size	Bare Nuclei	Bland Chromatin	Normal Nucleoli	Mitoses	Class
1057013	8	4	5	1	2	4	7	3	1	Malignant
1096800	6	6	6	9	6	2	7	8	1	Benign
1183246	1	1	1	1	1	2	2	1	1	Benign
1184840	1	1	3	1	2	2	2	1	1	Benign
1193683	1	1	2	1	3	3	1	1	1	Benign

1193698	1	1	2	1	3	3	1	1	1	?
---------	---	---	---	---	---	---	---	---	---	---

<https://www.kaggle.com/datasets/yasserh/breast-cancer-dataset>

## Example Tree



### Each node shows

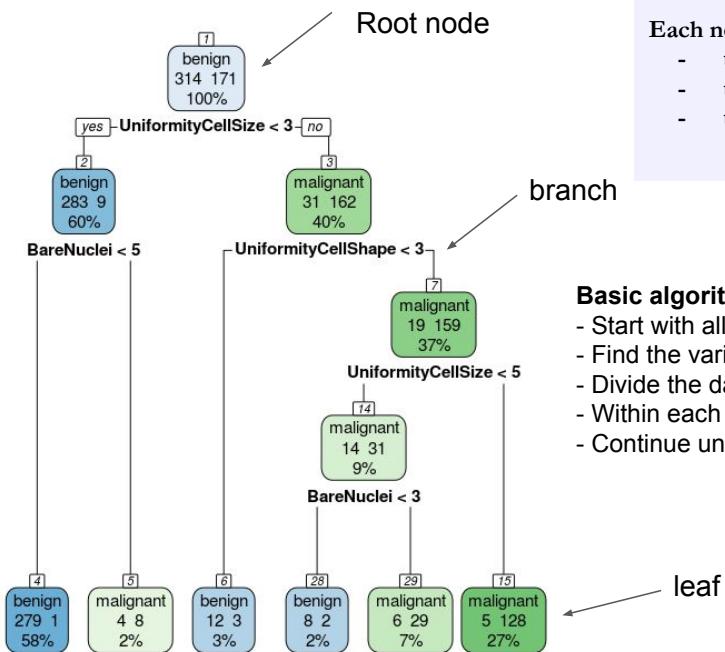
- the predicted class (benign or malignant)
- the number of observations per class
- the percentage of observations with respect to the total

### Basic algorithm

- Start with all variables in one group
- Find the variable/split that best separates the outcomes
- Divide the data into two groups ("leaves") on that split ("node")
- Within each split, find the best variable/split that separates the outcomes
- Continue until the groups are too small or sufficiently "pure"

Clump Thickness	Uniformity of Cell Size	Uniformity of Cell Shape	Marginal Adhesion	Single Epithelial Cell Size	Bare Nuclei	Bland Chromatin	Normal Nucleoli	Mitoses
-----------------	-------------------------	--------------------------	-------------------	-----------------------------	-------------	-----------------	-----------------	---------

## Example Tree



Root node

branch

leaf

Each node shows

- the predicted class (benign or malignant)
- the number of observations per class
- the percentage of observations with respect to the total

### Basic algorithm

- Start with all variables in one group
- Find the variable/split that best separates the outcomes
- Divide the data into two groups ("leaves") on that split ("node")
- Within each split, find the best variable/split that separates the outcomes
- Continue until the groups are too small or sufficiently "pure"

# Measures of impurity

It is the probability of being wrong if we predict the majority class.

Misclassification Error:

$$\hat{p}_{mk} = \frac{1}{N_m} \sum_{x_i \text{ in Leaf } m} 1(y_i = k)$$

Number of times that class k appears in leaf m

N is the number of total objects in leaf m

$$1 - \hat{p}_{mk(m)}; k(m) = \text{most; common; k}$$

- 0 = perfect purity
- 0.5 = no purity

Gini index:

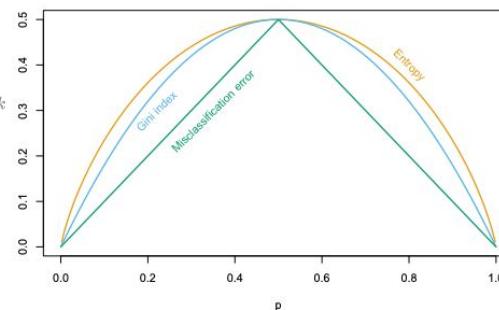
how likely it is that two randomly chosen observations from a node belong to different classes,

- 0 = perfect purity
- 0.5 = no purity

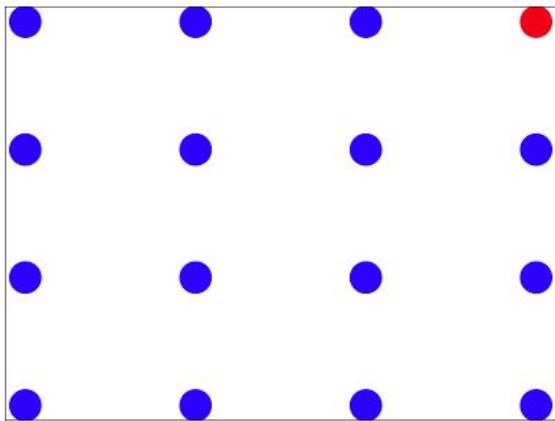
For 2 classes

$$\sum_{k \neq k'} \hat{p}_{mk} \times \hat{p}_{mk'} = \sum_{k=1}^K \hat{p}_{mk} (1 - \hat{p}_{mk}) = 1 - \sum_{k=1}^K p_{mk}^2$$

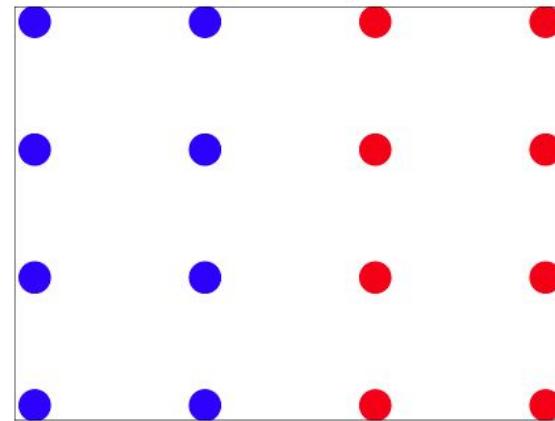
the sum of the squared probabilities of each class



## Measures of impurity



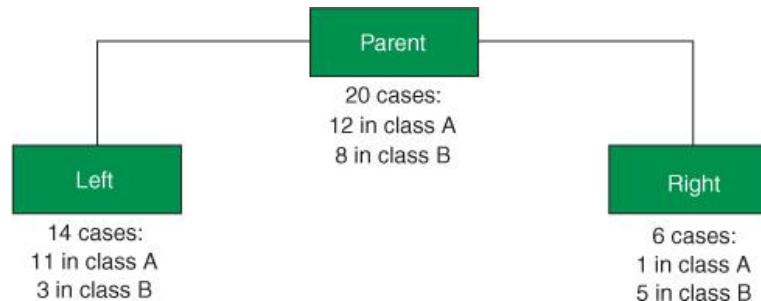
- Misclassification:  $1/16 = 0.06$
- Gini:  $1 - [(1/16)^2 + (15/16)^2] = 0.12$



- Misclassification:  $8/16 = 0.5$
- Gini:  $1 - [(8/16)^2 + (8/16)^2] = 0.5$

## Predicting with trees

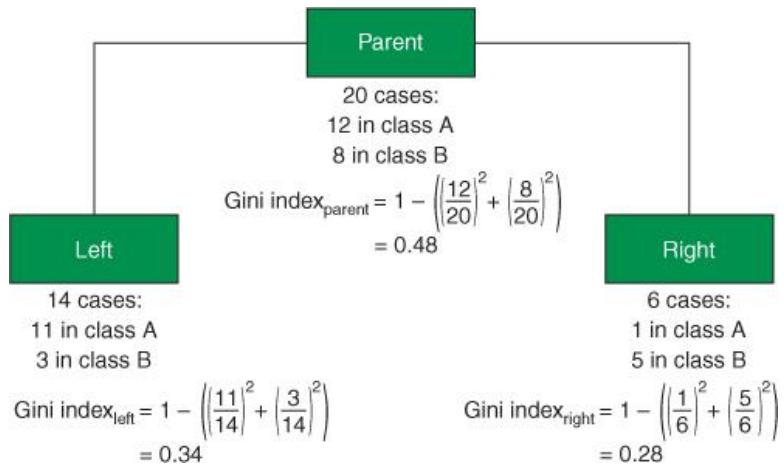
DMI



the Gini gain of this split is the difference between the Gini index of the parent node and the Gini index of the split.

$$\text{Gini index} = 1 - (p(A)^2 + p(B)^2)$$

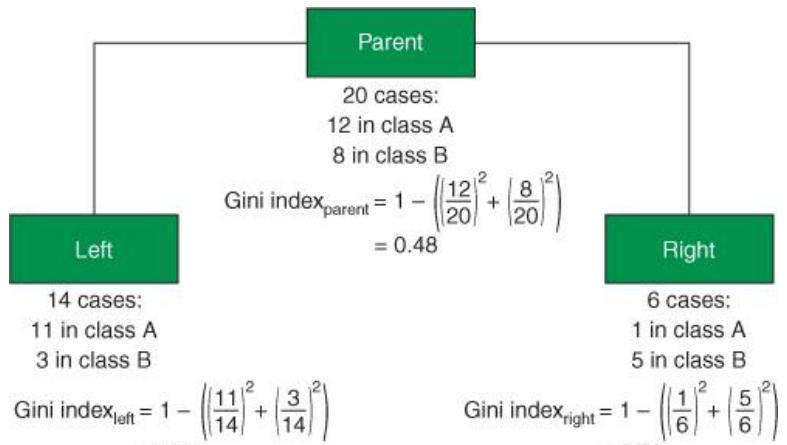
## Predicting with trees



$$\text{Gini index} = 1 - (p(A)^2 + p(B)^2)$$

The Gini gain at a particular node is calculated for each predictor variable, and the predictor that generates the largest Gini gain is used to split that node. This process is repeated for every node as the tree grows

## Predicting with trees



$$\text{Gini index}_{\text{split}} = p(\text{left}) \times \text{Gini index}_{\text{left}} + p(\text{right}) \times \text{Gini index}_{\text{right}}$$

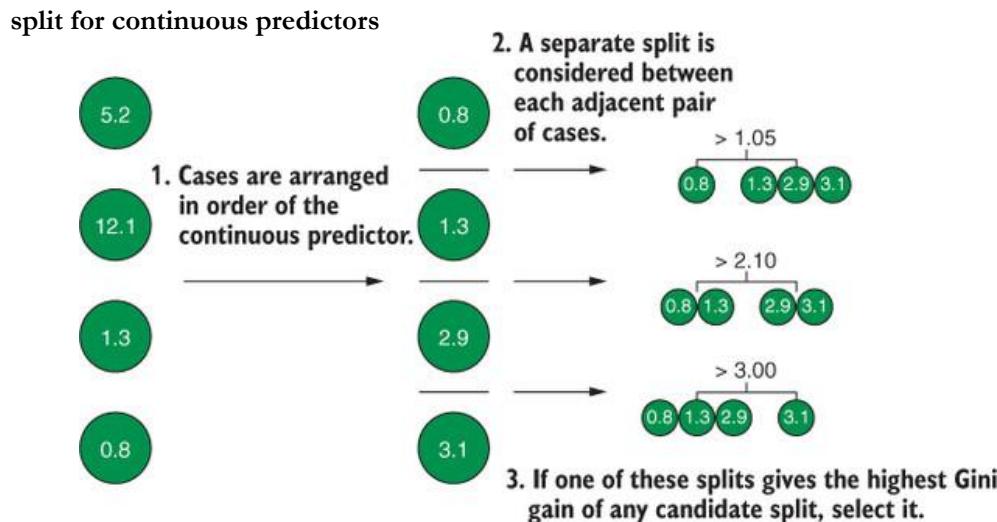
$$\text{Gini index}_{\text{split}} = \frac{14}{20} \times 0.34 + \frac{6}{20} \times 0.28 = 0.32$$

$$\text{Gini gain} = 0.48 - 0.32 = 0.16$$

$$\text{Gini index} = 1 - (p(A)^2 + p(B)^2)$$

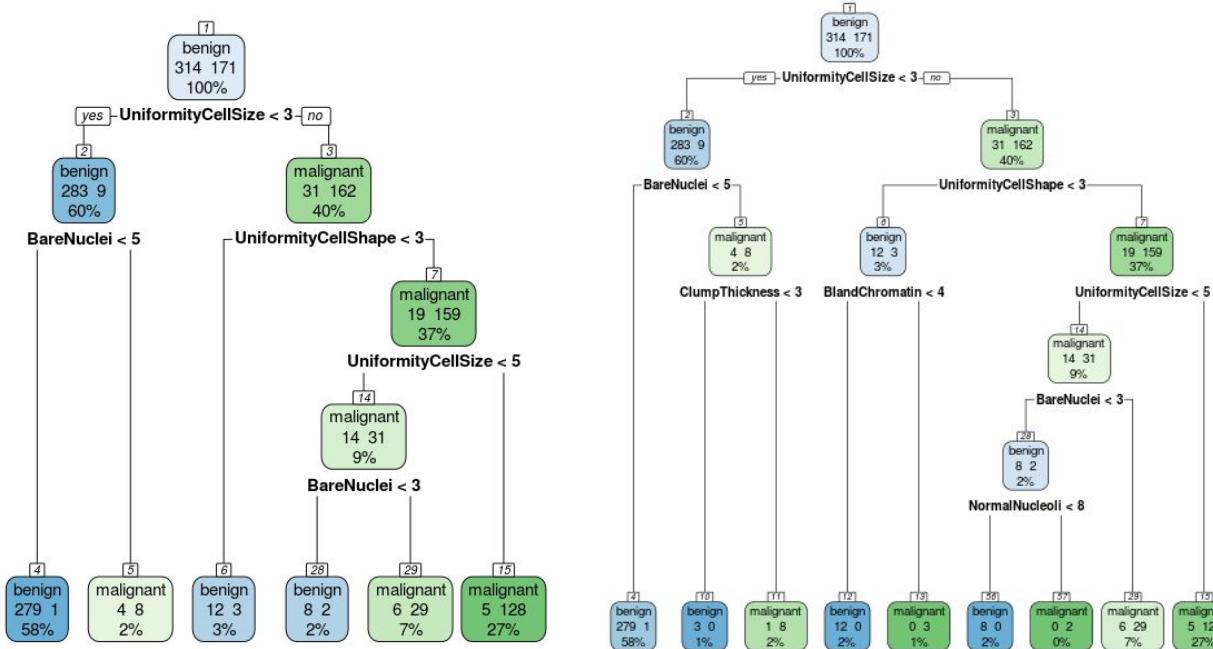
The Gini gain at a particular node is calculated for each predictor variable, and the predictor that generates the largest Gini gain is used to split that node. This process is repeated for every node as the tree grows

## Predicting with trees



# How deep?

**DMI**



## Model Selection for Decision Trees

### Pre-Pruning (Early Stopping Rule)

- Stop the algorithm before it becomes a fully-grown tree
- Typical stopping conditions for a node:
  - Stop if all instances belong to the same class
  - Stop if all the attribute values are the same
- More restrictive conditions:
  - Stop if number of instances is less than some user-specified threshold
  - Stop if expanding the current node does not improve impurity measures (e.g., Gini or information gain).
  - Stop if estimated generalization error falls below certain threshold

# Model Selection for Decision Trees



## Post-Pruning

- Grow decision tree to its entirety
- Subtree replacement

Trim the nodes of the decision tree in a bottom-up fashion

If generalization error improves after trimming, replace sub-tree by a leaf node

Class label of leaf node is determined from majority class of instances in the sub-tree

# Decision trees

## Strengths

- ✓ Easy to Understand.
- ✓ Useful in Data exploration. Decision tree is one of the fastest way to identify most significant variables and relation between two or more variables.
- ✓ Less data cleaning required. It requires less data cleaning compared to some other modeling techniques. It is not influenced by outliers and missing values to a fair degree.
- ✓ Data type is not a constraint. It can handle both numerical and categorical variables.
- ✓ Decision tree is considered to be a non-parametric method. This means that decision trees have no assumptions about the space distribution and the classifier structure.

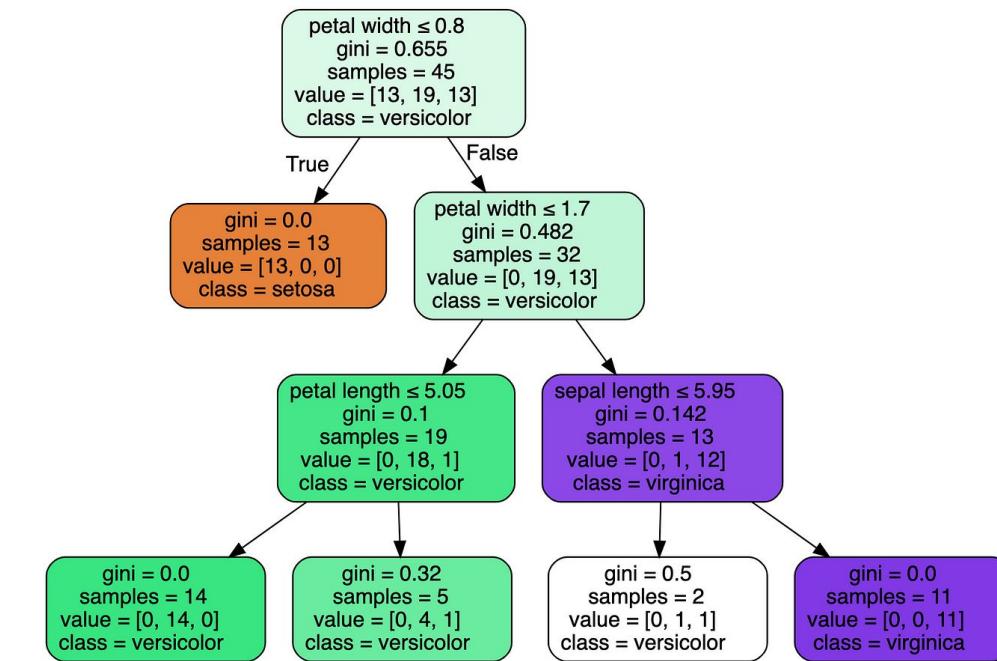
# Decision trees



## Weaknesses

- ✓ Overfitting: This problem gets solved by setting constraints on model parameters and pruning
- ✓ trees can be very non-robust: a small change in the data can cause a large change in the final estimated tree.
- ✓ individual decision trees generally do not often achieve state-of-the-art predictive accuracy.

## Multiclass decision tree





## How do we know how good our model is?

## An example: the breast cancer dataset



Sample number	True Class	Predicted Class
1057013	Malignant	Malignant
1096800	Benign	Benign
1183246	Benign	Malignant
1184840	Benign	Benign
1193683	Benign	Benign

## Metrics for Performance Evaluation



- Confusion Matrix:

		PREDICTED CLASS	
		Class=Yes	Class>No
ACTUAL CLASS	Class=Yes	a (TP)	b (FN)
	Class>No	c (FP)	d (TN)

- a: TP (true positive)
- b: FN (false negative)
- c: FP (false positive)
- d: TN (true negative)

$$\text{Accuracy} = \frac{a + d}{a + b + c + d} = \frac{TP + TN}{TP + TN + FP + FN}$$

## Limitation of Accuracy

- Consider a 2-class problem
  - Number of Class 0 examples = 9990
  - Number of Class 1 examples = 10
- If model predicts everything to be class 0,  
accuracy is 99.9 %

		Predicted Class	
		disease	Non-disease
Actual Class	disease	0	10
	Non-disease	0	9990

Accuracy is misleading because model does not detect any class 1 example

## Cost-Sensitive Measures

**DMI**

$$\text{Precision (p)} = \frac{a}{a + c}$$

$$\text{Recall (r)} = \frac{a}{a + b}$$

$$\text{F - measure (F)} = \frac{2rp}{r + p} = \frac{2a}{2a + b + c}$$

$$\text{Precision} = \frac{tp}{tp + fp}$$

$$\text{Recall} = \frac{tp}{tp + fn}$$

		PREDICTED CLASS	
		Class=Yes	Class>No
ACTUAL CLASS	Class=Yes	a (TP)	b (FN)
	Class>No	c (FP)	d (TN)

## Receiver Operating Characteristic (ROC) curve



In binary classification you are predicting one of two categories

Alive/dead

Disease/Healthy

But your predictions are often quantitative:

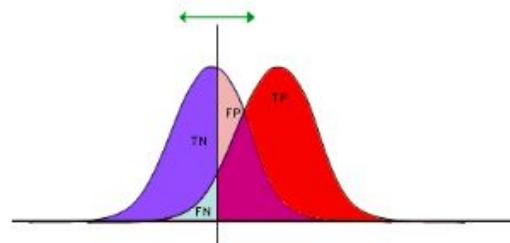
Probability of being alive

The cutoff you choose gives different results

## Receiver Operating Characteristic (ROC) curve

DMI

- AUC = 0.5: random guessing
- AUC = 1: perfect classifier
- In general AUC of above 0.8 considered "good"

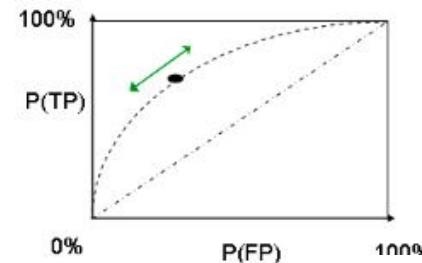


TP	FP
FN	TN
1	1

True positive rate (recall, sensitivity)

$$TPR = \frac{TP}{TP + FN}$$

AUC: Area under the curve



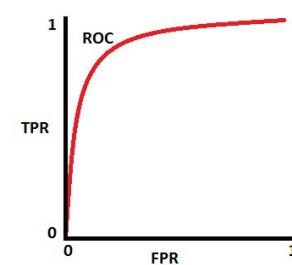
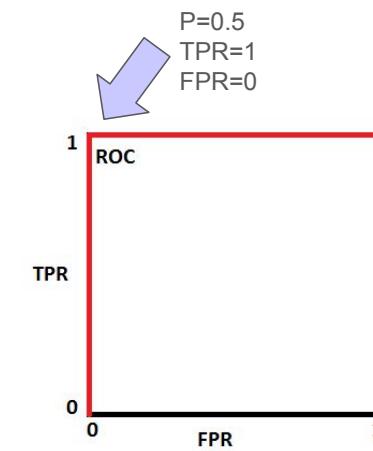
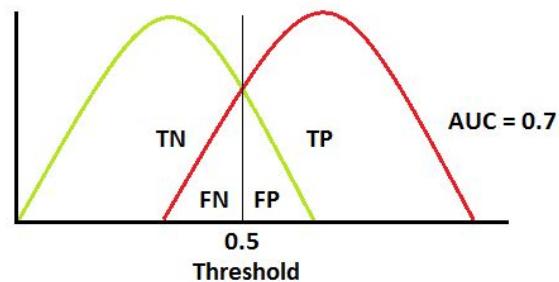
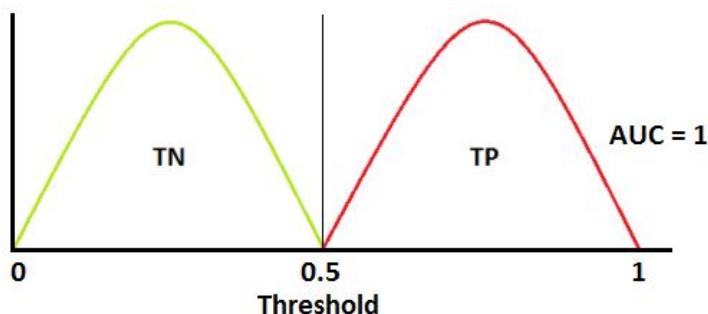
False Positive Rate

$$FPR = \frac{FP}{FP + TN}$$

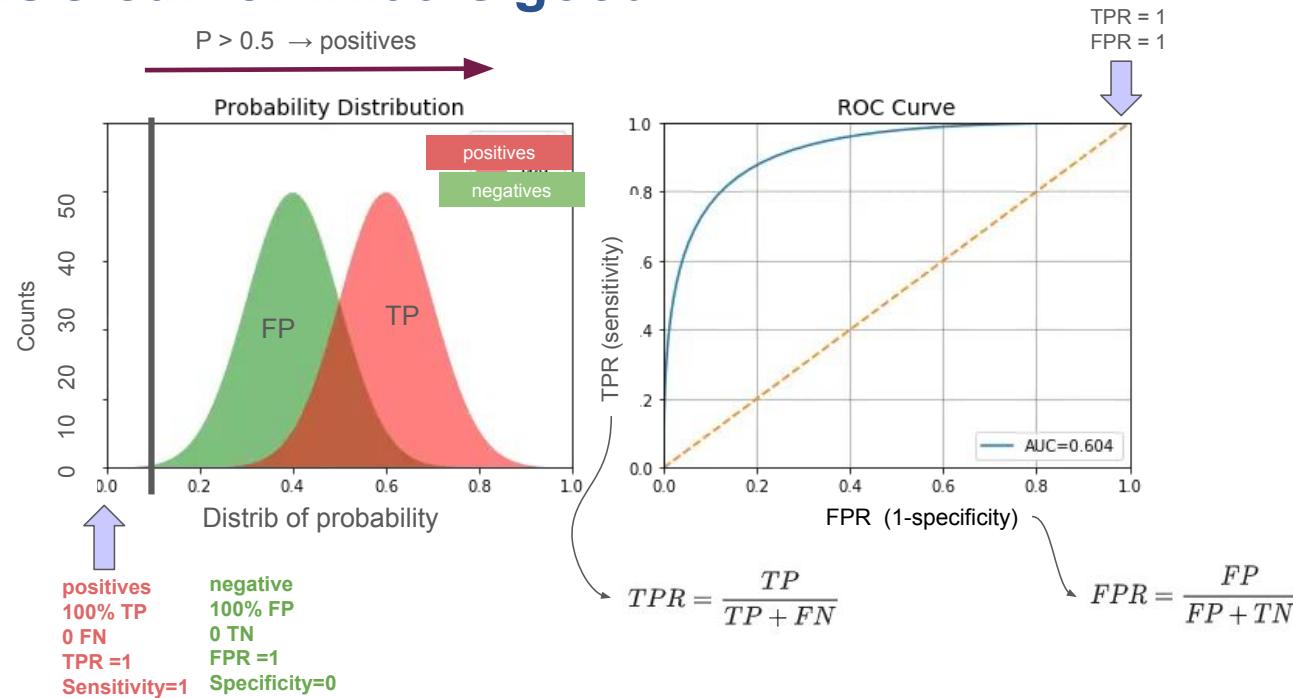
probability of false alarm, FPR= 1-specificity

## ROC curve: what is good?

DMI

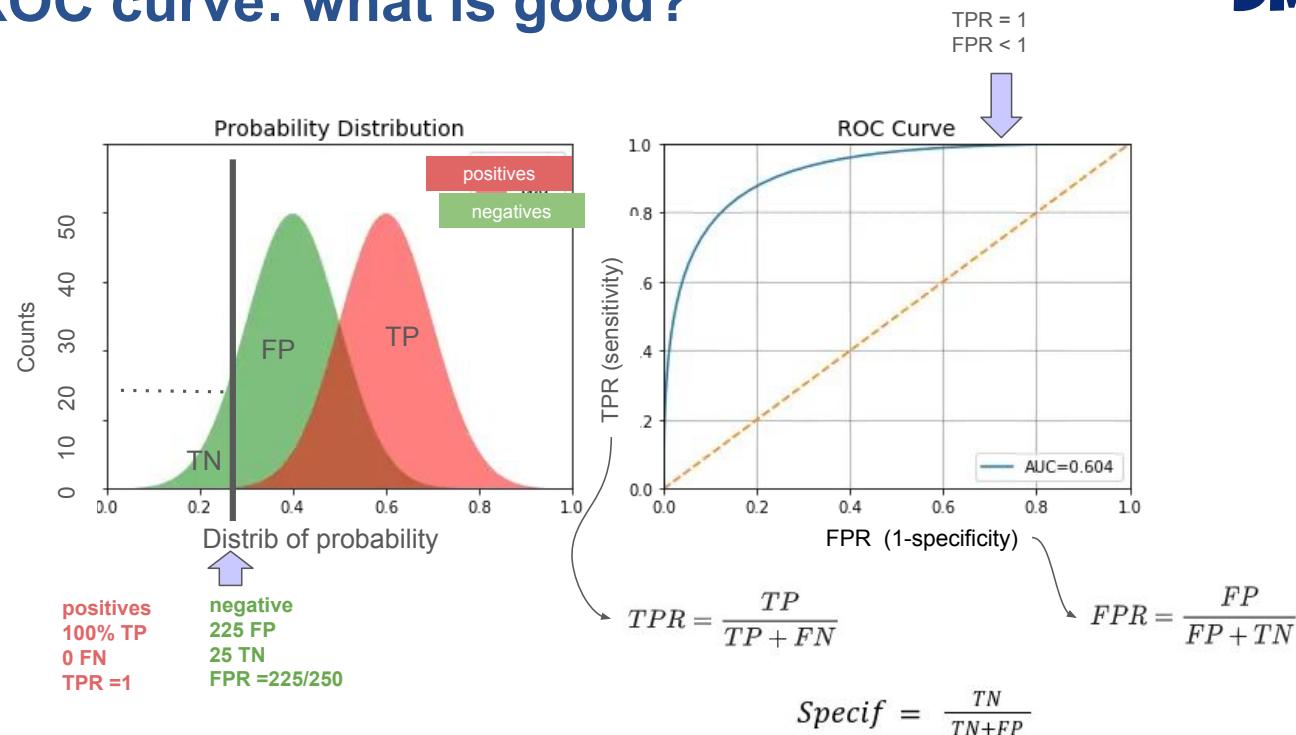


## ROC curve: what is good?



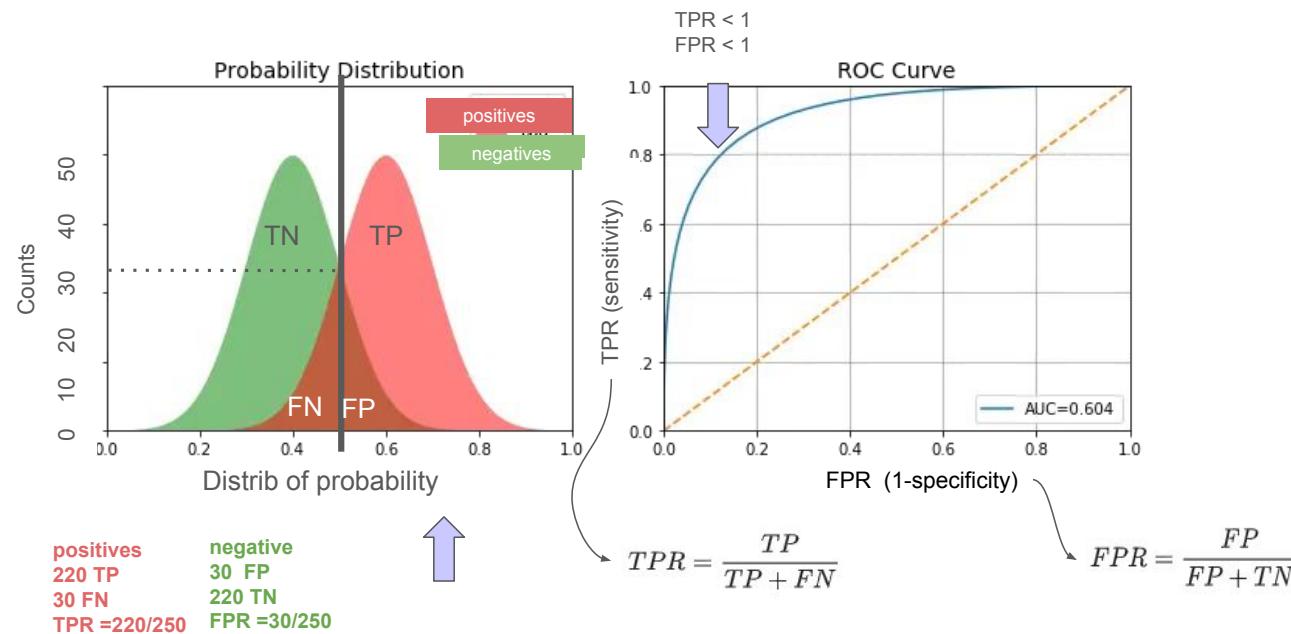
# ROC curve: what is good?

**DMI**



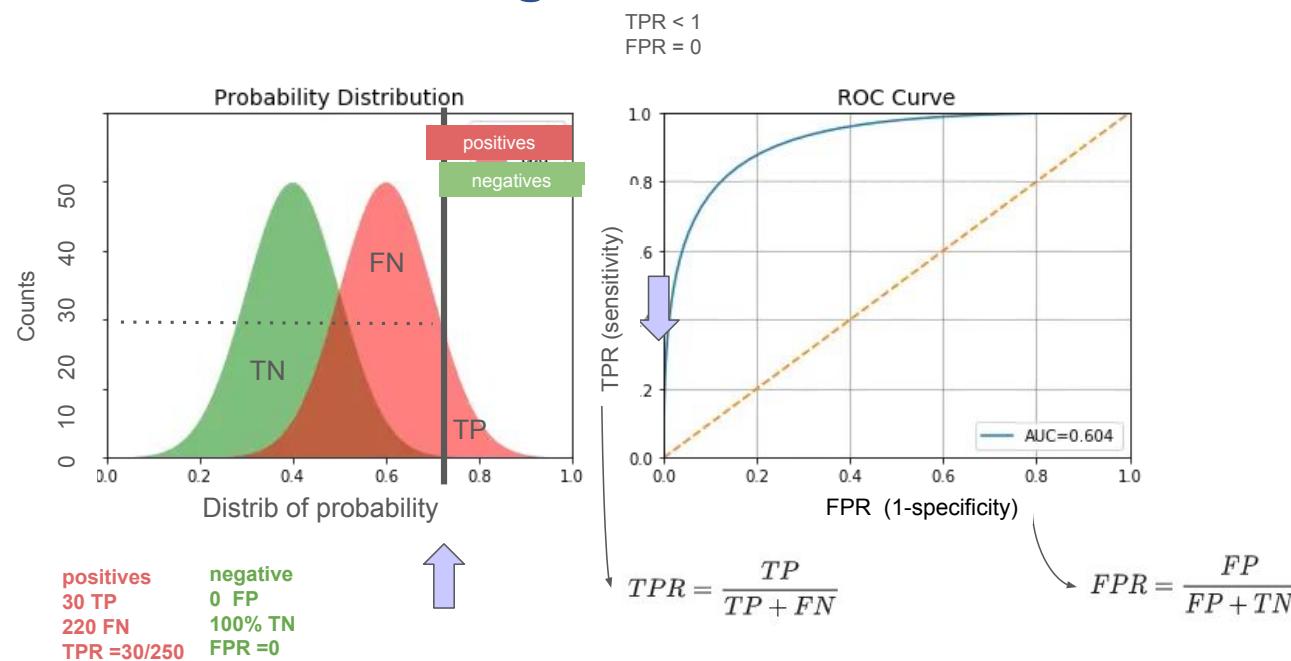
# ROC curve: what is good?

DMI



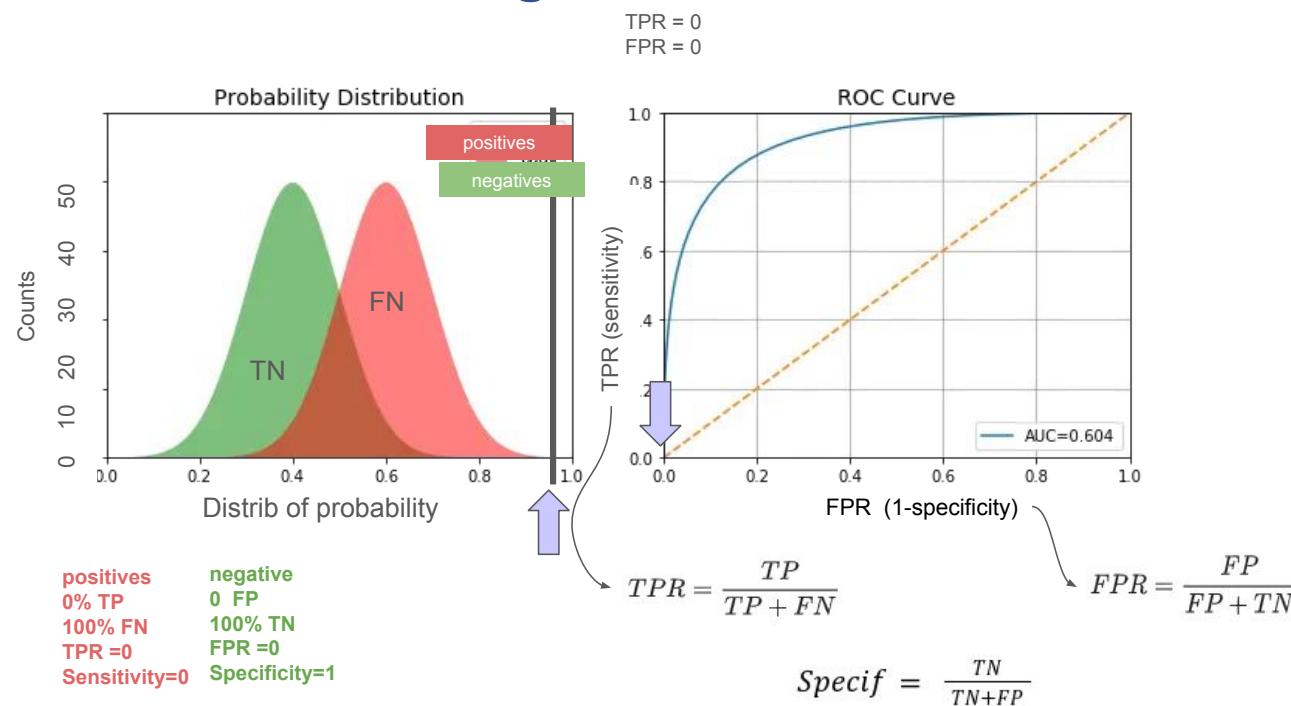
# ROC curve: what is good?

**DMI**

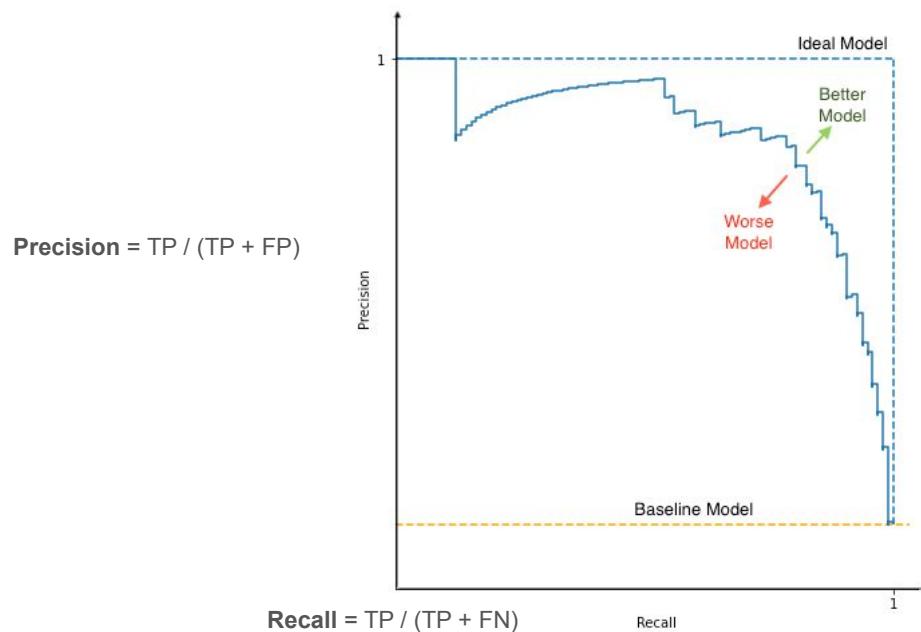


# ROC curve: what is good?

**DMI**



## Precision-Recall curves



Precision-Recall curves are often considered more informative than ROC curves when dealing with unbalanced classes. This is because Precision and Recall focus on the performance of a classifier in the positive class, which can be crucial when one class is significantly smaller than the other.

## The tidymodels framework

- ✓ **parsnip** for model definition
- ✓ **recipes** for data preprocessing and feature engineering
- ✓ **rsample** to resample data (useful for cross-validation)
- ✓ **yardstick** to evaluate model performance
- ✓ **dials** to define tuning parameters of your models
- ✓ **tune** for model tuning
- ✓ **workflows** which allows you to bundle everything together



## **Hands on session**



Go to this [page](#) and accept the assignment

## Bibliography



Statistics versus machine learning. Nat Methods 15, 233–234 (2018).

<https://doi.org/10.1038/nmeth.4642>

**Elements of Statistical Learning**, Data Mining, Inference, and Prediction

Trevor Hastie, Robert Tibshirani, Jerome Friedman

**An Introduction to Statistical Learning**, with Applications in R

Gareth James, Daniela Witten, Trevor Hastie and Robert Tibshirani